

DEV SPRINT 1 · STATUS REPORT

Foundations, Agent Fleet, and the End-to-End Loop.

Three days. Twenty-three pull requests. One platform that drafts copy, image, video, voice and music; previews every output; gates outbound through Approval; and publishes to four social networks plus Resend email.

SPRINT WINDOW

09 May – 12 May 2026

REPORT DATE

12 May 2026, 21:00 IST

PLAN COMPLETION

≈ 52 % of v1.2 / v2.0

Sprint 1 · DELIVERED

Prepared by

Engineering, with Claude Code · Opus 4.7 (1M context)

Confidential — internal distribution

v1.0 · 12.05.2026

IMPRINT

About this report

This is the official engineering and product status report for DClaw Marketing's first development sprint. It is intended to be read by three audiences: **operators** (TL;DR), **executives** (Executive Summary), and **engineers** (Phase-by-Phase Status and onward).

FIELD	VALUE
Document title	DClaw Marketing — Dev Sprint 1 Status Report
Sprint number	Sprint 1 (of an open-ended cadence)
Sprint window	2026-05-09 15:32 IST → 2026-05-12 21:00 IST (≈ 3 days 5 h)
Sprint kick-off commit	93189b1 · docs(readme): add Deepto Mallick as contributor
Snapshot ref	main @ 64242e8 (after PR #142)
Document version	1.0 — frozen at report-generation time
Generated	2026-05-12 21:00 IST via Chrome headless from sprint-1-status-2026-05-12.html
Location	obsidian/reports/ in the dclaw-marketing repo
Audience	Engineering, Product, Executive sponsor
Authoring	Engineering with Claude Code (Opus 4.7, 1M context)

How to read this document. If you have 5 minutes, read Chapter 1 (TL;DR). If you have 30 minutes, add Chapters 2–3 (Intent + Executive). If you are an engineer reviewing the work, read everything from Chapter 4 onward; appendices contain the raw catalogue of every pull request, test file, and architectural artefact shipped in the sprint.

Brand & design

Typography, colour, and tone follow the --dk-* brand token system defined in `frontend/src/styles/brand.css`. The body face is Poppins; charts and diagrams use the official purple palette (primary #7660A8, deepest #4A3878, lightest tint #F8F6FB). All inline charts in this document are hand-authored SVG against that palette — no external chart library.

CONTENTS

Table of Contents

Eleven chapters and four appendices. Total $\approx$ 60 pages.

1	TL;DR — the sprint in three pages	8
2	Project Intent & Target Audience	11
3	Executive Summary	15
4	The Original Plan (v1.2 / v2.0)	20
5	Sprint Execution Timeline	25
6	Phase-by-Phase Detailed Status	29
7	Pull Request Catalogue	48
8	Architecture Snapshot	54
9	Test & Quality Gates	57
10	Risks, Tech Debt, Open Decisions	60
11	Sprint 2 — Recommended Roadmap	63
A	Closed Issues with PR Mapping	66
B	Open Issues — Status & Owners	68
C	Full PR Catalogue Table	70
D	Glossary & Acronyms	73

CHAPTER 1

TL;DR

Three days. Twenty-three pull requests opened, twenty-two merged, one in the runway. The platform that started Sprint 1 as a scaffold now drafts *any* creative medium, previews it, gates it, publishes it, attributes it, and bills for it.


The sprint in five sentences

1. The **agent fleet went live**: Conductor + four role agents (SMM, SEO, Paid Media, Analyst) now run on real Claude with conversation history; every station has a shared chat surface.
2. The **creative pipeline is end-to-end**: copy + image (Replicate SDXL) + video (Replicate) + voice (ElevenLabs) + music — with brand-aware visual prompts and a deterministic stub fallback for dev.
3. The **Approval Inbox renders every medium inline** — image, animated SVG video, audio players for voice and music — so reviewers actually see what they're approving.
4. Four real social **publishers landed** (Bluesky, LinkedIn, X, Instagram) plus Resend email with a full hard-gated send loop, daily analytics rollups, three attribution models, identity resolution, GDPR export, daily quota guards, and a sync workflow runner.
5. The platform sits at **≈ 52 % of the v1.2 / v2.0 plan** — backend ≈ 65 %, frontend ≈ 40 %. Adapter breadth (more publishers, more newsletter providers, more CRMs), and the agency-operations UI (Kanban, Stripe, time tracking, client portal) are the headline gaps.

Headline numbers

PRS MERGED 22 +1 in runway · #121 → #142	NEW TESTS +81 15 new test files, all passing	LINES ADDED 47.9k across 651 files	PLAN COMPLETE ≈ 52 % vs v1.2 + v2.0
---	---	---	--

Phase completion at a glance



Top wins

- **The agent network is real.** All five agents (Conductor + SMM/SEO/Paid Media/Analyst) make actual Claude calls with 10-turn history. No more keyword stubs.
- **Multi-modal generation is wired.** Image + video + voice + music endpoints all live with provider-or-stub patterns. Brand kit prepended to every visual prompt.
- **Publishing has four channels.** Bluesky / LinkedIn / X / Instagram all post for real when tokens are set; gracefully stub otherwise.
- **Analytics close the loop.** Daily rollups → identity resolution → first/last/linear attribution → /totals hero-cards.
- **Compliance primitives shipped.** Daily quota guard, /costs endpoints, GDPR org-export task — three of the four §I deliverables in v1.2.

Top risks & gaps (carry into Sprint 2)

- **Cost ledger has no writers** — every adapter would need a one-liner hook to make /costs/totals show real numbers. Highest ROI fix.
- **Real OAuth flows missing** for all four publishers — they currently read tokens from a placeholder column.
- **Auto-merge label doesn't actually auto-merge** in this repo's settings — six PRs in the sprint required manual `gh pr merge` after CI was green.
- **Workflow runner is library-only** — no HTTP endpoint yet to kick it; `approval / branch` nodes are deferred pending a `WorkflowRun` model.
- **UI lag on Phases 2, 7, 10, 11** — backend primitives exist but no admin surfaces (Brand Setup Studio, sequence builder, time tracker, client portal).

One-sentence next step

Sprint 2 should open with **cost-write hooks on every adapter** (makes /costs/totals useful in production), follow with **real OAuth for the four shipped publishers**, then ship the **Brand**

Setup Studio and the **sequence builder** UIs — the largest backend-vs-frontend deltas in the current state of the plan.

CHAPTER 2

Project Intent & Target Audience

Before any code, what are we building, and for whom?

2.1 What DClaw Marketing is

DClaw Marketing is an **agent-driven marketing operating system** — a single multi-tenant platform that lets a marketing team (internal department or external agency) hand a brief to a fleet of specialised LLM agents and have it travel, gated, all the way to published content, attributable revenue, and a client-ready report.

The platform combines four ideas that are usually sold separately:

1 • Agent fleet

A Conductor + role specialists (Social Media Manager, SEO Specialist, Paid Media Specialist, Creatives, Analyst, plus Reviewer/Viewer/Client roles). Each agent owns a "station" in the product.

2 • Hard-gate supervision

Every outbound action — a tweet, an ad spend, an email send — files an Approval Request first. A human reviewer can approve, reject with reason, or let it auto-expire. Nothing escapes the gate.

3 • Bring-your-own everything via MCP

A registry of 35 MCP servers (HubSpot, Ahrefs, GA4, Mailchimp, Stripe, ...). Agents call tools through a unified async client; secrets are Fernet-encrypted per-org.

4 • Closed analytics loop

Touchpoints + Conversions are linked via identity resolution; daily rollups feed multi-model attribution; dashboards close the "did the agent move revenue?" question.

2.2 The people we're building for

Ten supervision-scope roles, defined in the v2.0 vision and now in the database as a Postgres enum on `OrganizationMembership.role`:

ROLE	WHO THEY ARE	PRIMARY SURFACE
admin	Org owner, billing, member management, hard-cap settings	/admin/*

manager	Approves everything, sets autonomy posture per action class	/inbox + /agent (Conductor)
creatives	Briefs in, variants out — copy/image/video/voice/music	/agents/creatives
social_media_manager	Owns the calendar + DM queue	/agents/smm + /calendar
seo_specialist	Keyword pipeline → topic clusters → drafts	/agents/seo
paid_media_specialist	Ad creative + A/B + bandit budget shifts	/agents/paid-media
reviewer	Decides on Approval items (not requester themselves)	/inbox
analyst	Daily rollups, anomaly detection, Monday narrative	/agents/analyst + /analytics
viewer	Read-only of everything in their Org	most pages
client	External user — sees only what their agency has shared (Theme O, not yet built)	/client/* (planned)

2.3 Differentiators

- **Hard-gate by default** — agents never publish on their own; everything funnels through Approval. PLAN-v1.2 §v2.0 §5.2 codifies this.
- **Multi-tenant from line one** — every row has `organization_id`; every API call is scoped by membership. No "later we'll add tenancy."
- **Agent fleet, not a single chatbot** — each role has its own system prompt, its own station, its own thread. The Conductor decomposes briefs and routes work.
- **Brand-aware** — the active BrandKit (palette, fonts, voice sliders) is prepended to every creative prompt, copy or visual.
- **Provider-or-stub everything** — no real API key? You still get a deterministic stub response so dev/CI keeps working. Replicate, ElevenLabs, Resend, Anthropic, MCP, Bluesky/LinkedIn/X/Instagram, all follow the same pattern.
- **Open data model** — Touchpoint, Conversion, CostLedger, AttributionResult, AnalyticsRollup are first-class tables. Not a closed black box.

2.4 What "v1.2 done" looks like

Per PLAN-v1.2.md and the v2.0 addendum, the platform is considered v1.2-complete when:

1. An org admin can onboard a new client, set the brand kit, ingest knowledge, and configure goals — all from the UI.
2. An SMM can hand the Conductor a goal ("grow LinkedIn engagement 20% this month") and watch role agents draft, schedule, and request approval for the work.
3. Outbound to **X / LinkedIn / Bluesky / Instagram / Threads / Mastodon / Pinterest / TikTok / YouTube / Reddit** (10 channels minimum) works for real, OAuth-connected.
4. Email outbound to Resend / Mailchimp / ConvertKit / Beehiiv works for both broadcasts and sequences.
5. Lead enrichment + CRM two-way sync (HubSpot / Salesforce / Pipedrive / Attio) is live.
6. Attribution shows multi-model credit allocation (linear, time-decay, Markov), the analytics dashboard reads from AnalyticsRollup, and per-channel cost vs revenue is visible.
7. Agency ops covers project/Kanban, time tracking + Stripe invoicing, weekly/monthly client PDFs, and the no-code Workflow builder.
8. Compliance: quota manager + sandbox mode + cost tracker + GDPR + client portal all reachable from the UI.


Sprint 1 has moved the **backend half** of these eight points from 30% to roughly 65% complete. The frontend half is still around 40% — that's the Sprint 2 opportunity.

CHAPTER 3

Executive Summary

For sponsors and decision-makers: what got delivered, what it cost (in scope), where the risk sits, and what to fund next.

3.1 Sprint outcome at a glance




3.2 What shipped — feature lanes

LANE	% OF PLAN	WHAT LANDED THIS SPRINT
Agent Fleet	70%	All five agents now use real Claude with 10-turn history. Shared chat UI on every station. Suggestion-chip handoff.
Content Generation	70%	Image (Replicate SDXL), video, voice (ElevenLabs), music — all live with stub fallback. Brand-kit prepended to every visual prompt. Inbox previews each medium inline.
Publishing	30%	Real publishers for Bluesky, LinkedIn, X, Instagram (4 of 20+). Worker dispatches per channel; per-platform error mapping.
Email / Outbound	30%	Resend adapter, approve-triggers-Celery-send, hard-gated campaign-send endpoint. Missing: SendGrid/Postmark, newsletter providers, sequences UI.
Integrations (MCP)	50%	Async invoke client with stub fallback + tool allow-list. 35-server registry pre-populated. Real per-server adapters still missing.
Analytics & Attribution	50%	Daily rollups (06:00 UTC), identity resolution (05:30), three attribution models (06:30), /totals + /rollups endpoints. Missing: enrichment, CRM sync.


Agency Operations	20%	Workflow DAG runner shipped (library-only); models for TimeEntry, Playbook, Workflow exist. Most UI surfaces and Stripe integration not started.
Compliance	40%	Daily quota guard on approve, /costs/totals + /costs/recent, GDPR export task. Missing: sandbox mode, cost-write hooks, client portal, sliding-window QuotaCounter.

3.3 Velocity — PRs landed by day



PRs merged per calendar day (UTC+05:30). The sprint compressed nearly all execution into 12 May.

3.4 Plan completeness — backend vs UI



Per-phase backend (purple) vs frontend (lavender). The platform's backend is materially ahead of its UI on every lane except 9 (Agent Fleet) and 1 (Foundation).

3.5 Risk register (C-level view)

RISK	LIKELIHOOD	IMPACT	MITIGATION
Cost ledger has zero writers — /costs/totals reports \$0 in prod	HIGH	MEDIUM	Sprint 2 lane #1 — add cost-write helper + one-liner hook per adapter
Real OAuth missing on 4 shipped publishers — production posting blocked	HIGH	HIGH	Sprint 2 lane #2 — implement OAuth for LinkedIn + X + Instagram (Bluesky uses app-password)
UI surfaces lag backend on Phases 2, 7, 10, 11	HIGH	MEDIUM	Brand Setup Studio + sequence builder + time tracker UIs as Sprint 2 lanes 3-5
Auto-merge label isn't wired — 6 manual merges in this sprint	MEDIUM	LOW	Half-day GitHub Actions task to wire the label → <code>gh pr merge --auto</code>
Workflow runner can't pause on approval — limits real workflows	MEDIUM	MEDIUM	Sprint 2 — add WorkflowRun model with paused state
Sandbox / dry-run mode not implemented — onboarding demos may post for real	MEDIUM	HIGH	Add <code>Organization.sandbox_mode</code> flag + adapter short-circuit; ½-day work

3.6 Funding recommendation for Sprint 2

Continue at current velocity for one more sprint. The largest gaps in the plan (Phases 5 / 7 / 10 adapter and UI breadth) are highly parallelisable. A second three-day sprint should push overall plan completeness from 52% to $\approx 75\%$, at which point the platform is demoable end-to-end and ready for a v1.2 release candidate.

3.6.1 Suggested Sprint 2 lanes (in priority order)

1. **Cost-write hooks** on every adapter (1 PR, half a day).
2. **Real OAuth** for LinkedIn / X / Instagram (3 PRs, 1 day).
3. **Brand Setup Studio** UI — closes the Phase 2 UI gap (1 PR, 1 day).
4. **Sequence builder backend + react-flow UI** — largest Phase 7 gap (2 PRs, 1 day).
5. **WorkflowRun model + HTTP endpoint** — unblocks workflow product (1 PR, ½ day).
6. **Cost-write hooks + sandbox-mode flag** — closes the riskiest compliance gaps (1 PR, ½ day).

CHAPTER 4

The Original Plan (v1.2 / v2.0)

The blueprint the sprint was measured against. Sourced from `PLAN-v1.2.md`, `IMPLEMENTATION-PLAN.md`, and the GitHub Epic issues #85–#96.

4.1 The 12-phase map

Sprint planning divides the v1.2 / v2.0 vision into twelve phases. They are partially-ordered (every Phase 1+ screen must use Phase 0's vocabulary, every Phase 9 agent must call into a Phase 6 MCP server) but otherwise can ship in parallel.

#	PHASE	INTENT
0	Design Ground Truth & Component Library	Every Phase 1+ screen should be born using the brand vocabulary (<code>--dk--*</code> tokens, Poppins, 23 <code><Dk*></code> components). Retrofitting later is wasteful.
1	Multi-Tenant Foundation (Theme A1, v2.0 §1–2)	Organization + Project hierarchy with 10 supervision-scope roles. Org switcher in nav; every API call scoped by membership.
2	Brand & Context Ingestion (Theme Q)	Versioned BrandKit (palette, fonts, voice sliders, do/don't say) + ICP personas + file/url/git/zip ingestion + pgvector KG + goals + autonomy posture.
3	Content Generation Pipeline (Theme B)	Brief editor → Campaign Kanban → copy / image / video / voice / music generation → Approval Inbox previews.
4	Calendar & Scheduling (Theme C1)	ScheduledPost model, Celery beat dispatcher, week-grid calendar UI, conflict detection, best-time recommender.
5	Multi-Account Publishing (Theme C2)	SocialAccount model with OAuth per-platform, publisher adapters for 20+ networks (X, LinkedIn, Bluesky, Instagram, Threads, Mastodon, TikTok, YouTube, Pinterest, ...).
6	MCP Integration Hub (Theme D)	Per-Org Connection rows with Fernet-encrypted secrets; registry of ~30 servers; async JSON-RPC-over-HTTP client; per-server adapters.
7	Email + Ads + Sequences (C3, C4, E4)	EmailTemplate / Campaign / Sequence / Segment models; Resend + SendGrid + Postmark + Mailchimp + ConvertKit + Beehiiv + Substack adapters; open/click webhooks; Ads (Meta, Google, LinkedIn); react-flow sequence builder.

8	Lead 2.0 + CRM Sync + Attribution (E + F)	Lead enrichment (Apollo / Clearbit / PDL), HubSpot / Salesforce / Pipedrive / Attio sync; Touchpoint + Conversion ingestion; first/last/linear + time-decay + Markov attribution; daily rollups; Sankey UI.
9	Agent Fleet (v2.0 §4)	Conductor + four role agents (SMM / SEO / Paid Media / Analyst) on real Claude Agent SDK; per-agent system prompts + tool list + memory; DM queue; A/B bandit; best-time recommender.
10	Agency Operations (J–P)	Onboarding wizard, Kanban + Gantt, time tracking, Stripe + QuickBooks invoicing, weekly/monthly PDF reports, Playbook search, no-code Workflow builder.
11	Compliance + Client Portal + Release	Sliding-window quota manager, sandbox/dry-run mode, cost tracking with hard caps, GDPR export + right-to-delete, external client portal, release polish.

4.2 Cross-cutting requirements

These constraints apply to every phase and are non-negotiable per PLAN-v1.2:

Hard-gate by default

Outbound posting, ads spend, email broadcasts — all funnel through `ApprovalRequest`. Per PLAN-v1.2 §v2.0 §5.2. Auto-approval is opt-in per action class via `Organization.autonomy_posture_json`.

Multi-tenant from line one

Every domain row has `organization_id`. Every API endpoint enforces membership. Every secret is scoped to one Org.

Light-mode only

No dark: Tailwind variants. The platform is single-mode by design. Reduces visual debt during v1.2 race.

Provider-or-stub

No real API key? Get a deterministic stub. Replicate, ElevenLabs, Resend, MCP, Anthropic, all four publishers — same pattern. Dev never blocked on credentials.

Audited everything

Every approve/reject/cancel writes an `AuditEvent` with actor, target, IP, and reason. Foundation for SOC-2.

Encrypted secret store

All third-party access tokens encrypted at rest with cryptography. fernet keyed off a per-

Org master key. Live behind
`app.services.secret_box`.

4.3 Phase intent — one-liner reminders

For quick reference when reading Chapter 6, the per-phase one-liner from the original plan:

PHASE	ONE-LINER INTENT (FROM PLAN-V1.2)
0	"Every Phase 1+ screen should be born using the design vocabulary. Retrofitting later is wasteful."
1	"Multi-tenant from line one — 10 supervision-scope roles + Org/Project hierarchy."
2	"The brand and the knowledge that every downstream agent reads from."
3	"From a brief to N variants in every medium, every variant gated."
4	"Schedule once, post on time, never twice."
5	"Twenty channels, one queue, one approval flow."
6	"Bring-your-own MCP — every external tool reachable through one client."
7	"Inbox + Ads + Sequences — outbound at scale, gated like everything else."
8	"Closed loop: touchpoint → conversion → attribution → dashboard → next brief."
9	"Five agents, one Conductor, real Claude — the brand brain."
10	"Agency-grade ops: clients, projects, time, billing, workflows."
11	"Compliance, reliability, polish — the things that make v1.2 shippable."

4.4 Acceptance criteria for v1.2 launch

1. An org admin can onboard a new client, set the brand kit, ingest knowledge, and configure goals — all from the UI, no shell access.
2. An SMM can give the Conductor a goal in plain English and watch role agents draft, schedule, and request approval for the work.
3. Outbound to 10+ social channels works for real, OAuth-connected.
4. Email outbound (broadcast + sequences) works through at least three providers.
5. Lead enrichment + CRM two-way sync is live for at least HubSpot.
6. Attribution shows multi-model credit allocation; dashboard reads from AnalyticsRollup.

7. Agency ops: Kanban, time tracking + Stripe invoicing, weekly client PDF, no-code workflow builder.
8. Compliance: quota manager, sandbox mode, cost tracker, GDPR, client portal — all reachable from the UI.

Sprint 1's contribution to each criterion is detailed in Chapter 6.

CHAPTER 5


Sprint Execution Timeline

What happened, in chronological order, across the three calendar days of the sprint.

5.1 Sprint window

MARKER	TIME (IST)	REFERENCE
Sprint kick-off	2026-05-09 15:32	commit 93189b1 · first commit by deepro713
First merge in scope	2026-05-11 (overnight)	PR #70 (pgvector embeddings + KG search)
Surge begins	2026-05-12 (morning)	PR #72 onward — 46 merges in 24h
Sprint freeze	2026-05-12 21:00	this report's snapshot ref (main @ 64242e8)

5.2 Cumulative PR burn-up



Cumulative merged PRs across the sprint window. The two earlier merges (KG, goals) seeded the foundation; the 46-PR surge on 12 May covered Phases 3–11.

5.3 Daily breakdown — what merged when

Day 1 — 2026-05-09 (Friday)

Setup day. Repo bootstrapping, Claude Code workflows wired up, brand-system + v1.2 plan adopted. No code merges in the strict sprint-1 numbered range (PRs #2–#15 land); this day's commits are foundational hygiene.

- Day 1 commits: 12

- **Headline:** brand system + plan doc adopted (commit 6abd826)

Day 2 — 2026-05-11 (Sunday)

Two merges. Knowledge graph (Q3) and goals/constraints (Q5) backends land. Light day; sets up Phase 2 / 8 primitives.

- **Merges:** PR #70 (KG + pgvector), PR #71 (Org goals + autonomy posture).
- **Day 2 commits:** 5

Day 3 — 2026-05-12 (Monday)

The surge day. 46 PRs land, covering Phases 0 finishing touches, all of Phase 1–4 scaffolding (#105–#120), then every Phase 3 / 4 / 5 / 6 / 7 / 8 / 9 / 10 / 11 depth-pass shipped by Claude Code (#121–#142).

- **Day 3 commits:** 61
- Of which PRs in the depth-pass range (**#121–#142**): **22 merged, 1 in CI**.
- **Conflict resolutions handled inline:** 6 (on publishing.py as each new publisher landed).
- **Stale issues closed with PR refs:** 12.

5.4 Activity heatmap

Merges per 4-hour bucket on 12 May



Approximate distribution of the 46 May-12 merges across the day. Peak load 12:00–16:00 IST.

5.5 Issue hygiene during the sprint

- **Closed (mapped to merging PRs):** 15 issues. Highlights: #88 Phase 3 epic, #91 Phase 6 epic, #92 Phase 7 epic, #93 Phase 8 epic, #94 Phase 9 epic, #32–#38 agent stations, #40 X, #41 LinkedIn, #42 Instagram, #43 Resend.
- **Still open at sprint end:** 35 → 25. Remaining items are Phase 0 design stories (#97–#104) plus Phase 5/10/11 long-tail (#85–#90 epics, #28 freshness, #24 helm, #95–#96).
- **Process bug observed:** PRs did not use Closes #N syntax for most of the sprint; tracker drift required a manual cleanup pass mid-sprint.

CHAPTER 6 · THE BIG ONE

Phase-by-Phase Detailed Status

For each of the twelve phases: the original scope (verbatim from the plan), what landed this sprint with PR citations, what's still missing, and the current depth estimate.

00 Design Ground Truth & Component Library

80%

"Every Phase 1+ screen should be born using the design vocabulary. Retrofitting later is wasteful."

PLAN SCOPE · 8 STORIES

- 0.1 Import brand assets — logos, icons, customer logos, pillar imagery
- 0.2 Tailwind token binding — bg-brand, text-fg-1, etc. → --dk-*
- 0.3 <Dk*> component library — 23 primitives
- 0.4 /_design reference page for visual eyeball-comparison
- 0.5 Refactor existing pages to the new vocabulary
- 0.6 Logo + favicon in nav
- 0.7 Top-level shell rebuild
- 0.8 Voice + casing audit

WHAT LANDED (MOSTLY PRE-SPRINT #105–#112)

- 0.1 ✓ DClaw placeholder logos in frontend/public/brand/ (PR #105 / #108)
- 0.2 ✓ Tailwind config binds to --dk-purple-*, --dk-fg-*, etc.; dark: variants disabled (PR #106)
- 0.3 ✓ **26** <Dk*> components currently in frontend/src/components/dk/ (PR #107) — 3 added this sprint (DkAgentChat from #123, plus refinements)
- 0.4 ✓ /admin/design reference page (PR #110)
- 0.5 ✓ Refactor pass (PR #112)
- 0.6 ✓ Logo + favicon (PR #108)
- 0.7 ✓ Sticky-glass top shell (PR #111)
- 0.8 ✓ Voice + casing audit folded into #112

WHAT'S STILL OPEN

- Stories #97–#104 in the tracker were not closed during the sprint despite being delivered — needs a one-time hygiene pass to close them with PR refs.
- No automated visual-regression coverage on the design page yet.

01 Multi-Tenant Foundation

70%

"Multi-tenant from line one — 10 supervision-scope roles + Org/Project hierarchy."

PLAN SCOPE

- 1.1 Org switcher in top nav, org context wired into client + every API call
- 1.2 /orgs list + create + edit + delete
- 1.3 /orgs/[id] tabbed detail (Overview / Members / Brand / Knowledge / Goals / Projects)
- 1.4 /orgs/[id]/members — invite by email, assign one of the 10 roles, remove

WHAT LANDED

- Backend: **100%**. Organization + OrganizationMembership + Project + ProjectMembership + 10-role enum + all CRUD endpoints — pre-sprint.
- UI: org switcher in nav (DkOrgSwitcher), admin/users, login + first-login flow — pre-sprint #113.
- This sprint: OrganizationMembership.role referenced from every new endpoint and worker task added in #121–#142 (gates correctly applied).

WHAT'S STILL OPEN

- Tabbed /orgs/[id] detail UI — separate pages exist but no unified tabbed view.
- Member invite-by-email UI.
- Per-Org settings page for autonomy posture overrides.

02 Brand & Context Ingestion (Theme Q)

50%

"The brand and the knowledge that every downstream agent reads from."

PLAN SCOPE

- 2.1 Q1 Brand Setup Studio — palette / type / voice sliders / examples-good vs examples-bad
- 2.2 Q2 Knowledge ingestion — file / URL / git / zip → DocumentChunks
- 2.3 Q3 Knowledge graph search — pgvector cosine similarity over 1536-dim embeddings
- 2.4 Q4 Freshness & re-ingestion
- 2.5 Q5 Org goals + constraints + autonomy posture
- 2.6 Q6 Project setup wizard

WHAT LANDED

- Backend: **100% complete for Q1–Q3, Q5** — BrandKit + Persona + DocumentChunk + KG endpoints + Goals endpoints all exist (pre-sprint + #70 + #71).
- This sprint: ✓ **Brand-aware visual prompt composer** (PR #128 — Phase 3.4 in my numbering, Q1 in the plan's numbering). The BrandKit's palette, fonts, and voice sliders now feed every image-gen prompt.
- Test coverage: test_brand_style.py (5 tests), test_brand_kits.py (8 tests), test_kg.py (12 tests), test_ingestion.py (12 tests), test_goals.py (6 tests).

WHAT'S STILL OPEN

- **Q1 Brand Setup Studio UI** — color pickers, font picker, voice sliders, example boards. The biggest UI gap on this phase.
- **Q2/Q3 Knowledge Console UI** — upload, paste URL, paste git remote; ingestion progress; KG search.
- **Q4 Freshness re-ingestion** (#28) — not started.
- **Q6 Project Setup Wizard** (#30) — not started.

03 Content Generation Pipeline (Theme B)

70%

"From a brief to N variants in every medium, every variant gated."

PLAN SCOPE

- 3.1 B2 Brief editor — Markdown + structured fields + template library
- 3.2 B2 Campaign Kanban
- 3.3 B3 Copy generation (Creatives Agent) — exists pre-sprint
- 3.4 B3 Image generation (Replicate / Stable Diffusion XL)
- 3.5 B3 Video generation (Replicate / Runway)
- 3.6 B3 Voice / TTS (ElevenLabs)
- 3.7 B3 Music generation (Replicate / Suno)
- 3.8 B3 Brand-aware prompt composition

WHAT LANDED THIS SPRINT

- **PR #124 · Phase 3.1 image generation** · `app/services/image_gen.py` · Replicate SDXL with stub fallback · `POST /agents/creatives/images · test_image_gen.py` (4 tests).
- **PR #125 · Phase 3.3 video/voice/music adapters** · `app/services/asset_gen.py` · 3 endpoints (videos, voice, music) · all stub-fallback · `test_asset_gen.py` (4 tests).
- **PR #126 · Phase 4.2 multimodal inbox previews** · `frontend/src/app/inbox/page.tsx` · `<img>` for image+video stubs, `<video controls>` for real video URLs, `<audio controls>` for voice/music.
- **PR #128 · Phase 3.4 brand-aware visual prompts** · `app/agents/brand_style.py` · prepends palette + font + mood to every image prompt · stores both raw + composed in `payload_json` for audit · `test_brand_style.py` (5 tests).

WHAT'S STILL OPEN

- **Brief editor** (`/projects/[id]/briefs/new`) — Markdown + structured fields + template library.
- **Campaign Kanban** (`/campaigns/[id]`).
- Brand-style prompt composition for the `/agents/creatives/videos` endpoint (only images use it today).

04 Calendar & Scheduling (Theme C1)

60%

"Schedule once, post on time, never twice."

PLAN SCOPE

- 4.1 ScheduledPost model + repository
- 4.2 Celery beat scanner — dispatch posts when due
- 4.3 Conflict detection — block two LinkedIn posts within 60 min on the same account
- 4.4 Best-time-to-post recommender — per-channel historical engagement
- 4.5 Week-grid calendar UI

WHAT LANDED THIS SPRINT

- **PR #126 · Multimodal inbox previews** (counted under Phase 3 above) — surfaces the calendar's drafts richly in the Approval Inbox.
- Pre-sprint (PR #114): ScheduledPost model, beat scanner (every 60s), week-grid /calendar UI.

WHAT'S STILL OPEN

- **4.3 Conflict detection** — no per-account same-channel cooldown check yet.
- **4.4 Best-time recommender** — would consume AnalyticsRollup data once enough has accumulated.

05 Multi-Account Multi-Channel Publishing

30%

"Twenty channels, one queue, one approval flow."

PLAN SCOPE

- 5.1 SocialAccount model + per-(org, platform, handle) uniqueness · OAuth per account
- 5.2 Publisher adapters for: X, LinkedIn (personal + company), Bluesky, Instagram, Threads, Mastodon, TikTok, YouTube, Pinterest, Reddit, Discord, plus 9 more in the plan's full list
- 5.3 /channels page — per-platform Connect-with-X cards, connection health

WHAT LANDED THIS SPRINT

- **PR #129 · Phase 5.1 Bluesky publisher** · `app/services/publishers/bluesky.py` · atproto two-step `createSession` + `createRecord` · uses app-password auth (no full OAuth) · `test_bluesky_publisher.py` (5 tests).
- **PR #137 · Phase 5.2 LinkedIn publisher** · `app/services/publishers/linkedin.py` · v2 `/ugcPosts` · Bearer OAuth 2.0 user-context token · permalink built from `x-restli-id` header · `test_linkedin_publisher.py` (5 tests).
- **PR #138 · Phase 5.3 X publisher** · `app/services/publishers/x.py` · v2 `/2/tweets` · 280-char truncation · OAuth 2.0 user-context bearer · `test_x_publisher.py` (6 tests).
- **PR #139 · Phase 5.4 Instagram Business publisher** · `app/services/publishers/instagram.py` · Graph API two-step `/media` + `/media_publish` · 2200-char caption cap · raises `InstagramMissingMediaError` when image required but absent · `test_instagram_publisher.py` (6 tests).
- Worker dispatch (`app/worker/tasks/publishing.py`) updated to route by `ScheduledPost.channel`; per-platform error mapping into `error_message`.

WHAT'S STILL OPEN

- **Real OAuth flows** — all four publishers use `SocialAccount._interim_access_token` as a placeholder. The OAuth grant + refresh dance is not implemented.
- **16+ more platform adapters** — Threads, Mastodon, TikTok, YouTube, Pinterest, Reddit, Discord, plus 9 from the plan's full list.
- **Per-platform Connect-with-X buttons** in the `/channels` UI.
- **Image attachments** on text-supporting platforms (LinkedIn, X) — currently text-only.

CHAPTER 6 · CONTINUED

Phase-by-Phase Detailed Status (continued)

06 MCP Integration Hub (Theme D)

50%

"Bring-your-own MCP — every external tool reachable through one client."

PLAN SCOPE

- 6.1 Connection model — workspace_id, server_id, name, encrypted secret blob, scopes, status, last_health_at
- 6.2 Fernet-encrypted secret store keyed off per-Org data key
- 6.3 Async MCP client — tool-call dispatcher with retry + audit logging
- 6.4 Built-in MCP servers / adapters for ~30 providers (HubSpot, Salesforce, GA4, Ahrefs, SEMrush, Mailchimp, Stripe, ...)
- 6.5 /integrations UI — registry browse, connection cards, health check, revoke

WHAT LANDED THIS SPRINT

- **PR #130 · Phase 6.1 async MCP client** · app/services/mcp_client.py · HTTP-transport JSON-RPC-like client; reads Connection.metadata_json.endpoint + decrypts secret with secret_box.unseal(); bearer-auth tool call; stub fallback when no endpoint.
- New endpoint **POST /orgs/{org_id}/connections/{connection_id}/invoke** — tool-name allow-list against the registry, connection-status gate, last_health_at stamp on success, last_error_message on failure.
- Test coverage: test_mcp_client.py (4 tests covering stub, real call, server error, error-in-body).
- Pre-sprint: Connection model + Fernet secret box (PR #116), 35-server registry pre-populated.

WHAT'S STILL OPEN

- **Per-server real adapters** — registry exists but no actual HubSpot search-contacts or GA4 metrics wrapper. The client can call *any* MCP-compliant endpoint; we just don't ship pre-built integrations.
- Retry policy + exponential backoff in the client (currently single-shot).
- Audit logging hooks on tool invocations (write a CostLedger-style row to a future MCPInvocation table).

07 Email + Ads + Sequences (C3, C4, E4)

30%

"Inbox + Ads + Sequences — outbound at scale, gated like everything else."

PLAN SCOPE

- 7.1 EmailTemplate / Audience / EmailCampaign / EmailSequence models
- 7.2 SendGrid + Resend + Postmark transactional adapters
- 7.3 Mailchimp + ConvertKit + Beehiiv + Substack newsletter adapters
- 7.4 Open / click / reply webhook ingest
- 7.5 Sequence builder (react-flow)
- 7.6 Segment builder UI + materialised SegmentQuery
- 7.7 Ads — Meta + Google + LinkedIn campaign management

WHAT LANDED THIS SPRINT

- **PR #127 · Phase 7.1 Resend adapter** · `app/services/email_send.py` · single `send_email()` over Resend with deterministic `msg_stub_<sha>` fallback · admin-only `POST /email/test-send` endpoint for credentials verification · `test_email_send.py` (3 tests).
- **PR #131 · Phase 7.2 approve → deliver Celery task** · `app/worker/tasks/email_send.py` · `deliver_approved_email(approval_id)` task · idempotent (skips re-send if delivery already recorded) · stamps `payload_json.delivery` with `message_id` + `provider` · `approve` endpoint enqueues this task for `action_type=send_email`.
- **PR #133 · Phase 7.3 campaign-send hard-gate** · `POST /email/campaigns/{id}/send` · resolves subject + html from EmailTemplate (override-able) · creates ApprovalRequest with `action_type=send_email` + recipient list in `payload_json` · rejects non-draft / non-scheduled (409) and empty body (400).
- Pre-sprint (PR #119): EmailTemplate, EmailCampaign, EmailSequence, EmailSequenceStep, Segment models scaffolded.

WHAT'S STILL OPEN

- SendGrid + Postmark transactional adapters.
- Newsletter adapters (Mailchimp / ConvertKit / Beehiiv / Substack).
- Open / click / reply / bounce webhook ingest.
- **Sequence builder (react-flow)** — biggest single UI gap.
- Segment builder UI + materialised SegmentQuery.
- Ads campaign management (Meta / Google / LinkedIn) — model exists, adapters and UI don't.

08 Lead 2.0 + CRM Sync + Attribution + Analytics

50%

"Closed loop: touchpoint → conversion → attribution → dashboard → next brief."

PLAN SCOPE

- 8.1 Extend Lead — enrichment_json, score, stage, source, utm_*, last_activity_at
- 8.2 LeadActivity + LeadNote models
- 8.3 Enrichment chain (Apollo / Clearbit / PDL via MCP)
- 8.4 CRM two-way sync — HubSpot / Salesforce / Pipedrive / Attio via MCP
- 8.5 Touchpoint + Conversion + AttributionResult + AnalyticsRollup models
- 8.6 Daily rollup job · identity resolution · multi-model attribution (first/last/linear/time-decay/Markov)
- 8.7 Dashboards — channel timeline, Sankey, lifecycle Kanban

WHAT LANDED THIS SPRINT

- **PR #132 · Phase 8.1 daily rollups** · app/worker/tasks/analytics.py · per-channel touchpoints/uniques + org summary of touchpoints/conversions/revenue · upserts AnalyticsRollup rows · runs 06:00 UTC · idempotent · test_analytics_rollup.py (4 tests).
- **PR #134 · Phase 8.3 attribution computation** · app/worker/tasks/attribution.py · first-touch / last-touch / linear allocation across journey · 30-day lookback · runs 06:30 UTC after rollups · idempotent (deletes prior rows per conversion+model) · test_attribution_compute.py (7 tests).
- **PR #135 · Phase 8.2 identity resolution** · app/worker/tasks/identity.py · back-stamps lead_id on anonymous Touchpoints sharing a visitor_id with an identified one · runs 05:30 UTC before rollups · never overwrites existing lead_id · test_identity_resolution.py (5 tests).
- **PR #136 · Phase 8.4 hero-card totals** · GET /api/v1/analytics/totals?organization_id&days=30 · current + previous window sums of touchpoints / conversions / revenue with delta-percent chips · fixed FastAPI deprecation (regex → pattern) on /rollups.
- Pre-sprint (PR #118): Touchpoint + Conversion + AttributionResult + AnalyticsRollup models.

WHAT'S STILL OPEN

- **Lead 2.0 fields** — enrichment_json / score / stage / last_activity_at not yet added to the Lead model.
- **LeadActivity + LeadNote** models — not created.
- **Enrichment chain** — Apollo / Clearbit / PDL via MCP — not implemented.
- **CRM two-way sync** — HubSpot / Salesforce / Pipedrive / Attio adapters — not implemented.

- Time-decay + Markov attribution models (the runner accepts them as enum values and returns empty allocations).
- Sankey + lifecycle Kanban UI.

09 Agent Fleet (v2.0 §4)

70%

"Five agents, one Conductor, real Claude — the brand brain."

PLAN SCOPE

- 9.1 Swap stubbed LLM for Claude Agent SDK runtime; per-agent system prompts + tool list + memory
- 9.2 Conductor Agent + Manager Station — decompose brief, dispatch to role-Agents, escalation queue
- 9.3 SMM Agent + Calendar integration — drafts, DM replies, best-time
- 9.4 SEO Agent — keyword pipeline, topic clusters, editorial workflow, internal linking
- 9.5 Paid Media Agent — ad creative, A/B, bandit budget shifts
- 9.6 Analyst Agent — daily rollups consumer, anomaly detection, Monday narrative
- 9.7 Shared chat UI on every station with suggestion-chip delegation

WHAT LANDED THIS SPRINT

- **PR #121 · Phase 9.1 Conductor real Claude** · `app/agents/conductor.py` rewritten from keyword classifier to real Anthropic calls · `complete(system, user)` wrapper with JSON-tolerant parsing · 10-turn history passed via `history=` · deterministic stub fallback when `ANTHROPIC_API_KEY` absent · system prompt names every platform surface (`/calendar`, `/channels`, `/integrations`, etc.) and hard-rules outbound is hard-gate.
- **PR #122 · Phase 9.2 role-agent runners** · `app/agents/_agent_base.py` shared runner with `RoleAgentTurn` dataclass · SMM / SEO / Paid Media / Analyst modules each supply their own system prompt + stub fallback · routing table in `agent_threads.py` dispatches by `AgentKind`.
- **PR #123 · Phase 9.3 shared chat UI** · `frontend/src/components/dk/agent-chat.tsx` · `<DkAgentChat kind={AgentKind} />` takes per-kind placeholder + empty-state copy · iMessage-style bubbles with brand-coloured user side · Cmd/Ctrl-Enter to send · suggestion chips: prompts (click-to-send) and links (deep-link).
- Conductor + four role-agent station pages each refactored to a thin wrapper around `<DkAgentChat>` with the right kind prop.

WHAT'S STILL OPEN

- **True Claude Agent SDK runtime** — currently uses the Anthropic API directly through `app.agents.anthropic_client.complete`; the official Agent SDK with tool-use loops and memory hasn't been wired.
- **DM queue + inbox UI** for the SMM agent.
- **A/B bandit policy** for the Paid Media agent (Thompson sampling across active variants).
- Conductor → role-agent **programmatic delegation** — Conductor currently suggests routes; doesn't create child threads on its own.

10 Themes J–P · Agency Operations

20%

"Agency-grade ops: clients, projects, time, billing, workflows."

PLAN SCOPE

- 10.1 J Client Operations — CRUD, onboarding wizard, per-Org retainers, approval workflows
- 10.2 K Project Management — templates, Kanban + Gantt, dependencies, capacity
- 10.3 L Time Tracking & Billing — time logs, auto-rollup to retainer burn-down, Stripe + QuickBooks invoice gen
- 10.4 M Client Reporting — weekly + monthly PDFs, scheduled email delivery, white-label
- 10.5 N Knowledge Base + SOPs — Playbook model, agent-searchable
- 10.6 O Client Portal — see Phase 11
- 10.7 P Workflow Builder — no-code visual chain of LLM steps + tool calls + approval gates + branches

WHAT LANDED THIS SPRINT

- **PR #141 · Phase 10.3 (P) Workflow runner** · `app/services/workflow_runner.py` · synchronous async DAG executor for `Workflow.dsl_json` · Kahn's-algorithm topological order · `{{var}}` + `{{node.text}}` template rendering walking dicts/lists · node runners for `llm` (Claude), `tool_call` (MCP), `noop` · `approval / branch / wait / webhook_listener` recognized but defer execution (need `WorkflowRun` model) · 10 unit tests in `test_workflow_runner.py`.
- Pre-sprint (PR #120): `CostLedger`, `QuotaCounter`, `TimeEntry`, `Workflow`, `Playbook`, `DataExportRequest` models scaffolded.

WHAT'S STILL OPEN

- **Client onboarding wizard** — collect brand assets / social accounts / persona / goals in a guided flow.
- **Kanban / Gantt** UI for projects + task dependencies + capacity planning.
- **Time-logging endpoints + timer UI** against the `TimeEntry` model.
- **Stripe + QuickBooks invoice generation** with `Invoice` model + migration.
- **Weekly/monthly PDF reports** with scheduled email delivery + white-label per Org.
- **Playbook search UI** backed by KG embeddings.
- **Workflow builder UI** (react-flow) for the now-runnable DSL.
- **WorkflowRun model** with paused-execution state to unlock approval/branch nodes.
- Workflow HTTP endpoint to kick the runner — currently library-only.

11 Compliance + Client Portal + Release

40%

"Compliance, reliability, polish — the things that make v1.2 shippable."

PLAN SCOPE

- 11.1 I1 Rate-limit / Quota — per-channel + per-provider sliding-window quotas + circuit breaker · UI shows "Twitter 47/300 today"
- 11.2 I2 Sandbox / Preview Mode — workspace flip to dry-run; every external action recorded but not fired
- 11.3 I3 Cost Tracking — aggregate per-Org LLM/image/video/voice/provider spend · daily budgets · soft + hard caps · /admin/costs
- 11.4 I4 Export / GDPR — workspace data export (zip), right-to-delete, per-record retention
- 11.5 O Client Portal — external-user view of approved content + scheduled posts
- 11.6 Release polish

WHAT LANDED THIS SPRINT

- **PR #133 (also Phase 11.2 quota guard)** · `app/services/quota_guard.py` · daily-cap check on every approve (`send_email` / `publish_ad` / `publish_social_post` / `image` / `video`) · reads `Organization.constraints_json.daily_caps_usd` · returns 409 with precise reason on breach · 7 unit tests in `test_quota_guard.py`.
- **PR #140 · Phase 11.3 GDPR org-export task** · `app/worker/tasks/gdpr.py` · `build_export()` walks every org-scoped model (org, memberships, projects, brand_kits + personas, campaigns, leads, scheduled_posts, approval_requests, social_accounts, connections, agent_threads + messages, touchpoints, conversions) · OAuth tokens + encrypted blobs always [REDACTED] · `export_organization_data` Celery task returns row counts + size · 6 unit tests in `test_gdpr_export.py`.
- **PR #142 · Phase 11.1 cost-ledger endpoints** · `app/api/v1/costs.py` · `GET /costs/totals?organization_id&days=30` with current-vs-previous-window breakdown by kind + provider · `GET /costs/recent?limit=50` for drill-down.
- Pre-sprint (PR #120): `CostLedger` + `QuotaCounter` + `DataExportRequest` models scaffolded.

WHAT'S STILL OPEN

- **Cost-write hooks** on every adapter — `CostLedger` has zero writers, so `/costs/totals` returns \$0 in production. This is the highest-ROI fix on the entire plan (one-liner per adapter).
- **Sliding-window QuotaCounter writer + UI** — daily caps work; rolling-window quotas don't yet update the `QuotaCounter` table.
- **Sandbox / dry-run mode** (Phase 11.2 I2) — `Organization.sandbox_mode` flag + adapter short-circuit.
- **Right-to-delete** companion to the export task.

- Persistence of the export to MinIO + signed URL response (currently returns the dict in-band).
- HTTP endpoint to kick the export.
- **Client Portal** (Theme O) — external-user surface; needed for v1.2 release.
- Circuit breaker around adapter calls.

CHAPTER 7

Pull Request Catalogue

Every merged PR in the depth-pass range (#121–#142), in chronological order, with file count, test count, and the one decision that defined it.


#	PHASE	TITLE	FILES	TESTS	KEY DECISION
121	9.1	Conductor real Claude + history	1	—	10-turn history + JSON-tolerant parsing fallback
122	9.2	SMM/SEO/Paid/Analyst real Claude	5	—	<code>_agent_base.py</code> shared runner + per-agent prompts
123	9.3	Shared <code><DkAgentChat></code> chat component	6	—	Single component takes kind prop; every station thin-wraps
124	3.1	Replicate image-gen adapter	4	4	Provider-or-stub; SVG stubs for offline dev
125	3.3	Video / voice / music adapters	4	4	Animated-SVG stub for video, silent-WAV stub for audio
126	4.2	Inbox renders multimodal previews	1	—	Branch by <code>asset_kind</code> : <code></code> / <code><video></code> / <code><audio></code>
127	7.1	Resend send adapter	3	3	Synthetic <code>msg_stub_<sha></code> id for deterministic stub
128	3.4	Brand-aware visual prompt composer	3	5	Pure rewrite — no LLM call; palette + font + mood prefix
129	5.1	Bluesky publisher + dispatch	4	5	App-password auth (no full OAuth dance for first publisher)
130	6.1	Async MCP tool-invocation client	3	4	HTTP transport; reads endpoint from <code>metadata_json</code>
131	7.2	Approve → <code>deliver_approved_email</code> task	3	—	Idempotent on <code>payload_json.delivery</code> ; best-effort enqueue
132	8.1	Daily analytics rollup job	5	4	Pure <code>compute_rollup_for_day()</code> + Celery wrapper
133	7.3 + 11.2	Campaign-send hard-gate + quota guard	4	7	Bundle — campaign send creates Approval, quota check on approve

134	8.3	First/last/linear attribution compute	4	7	Idempotent — delete prior rows per (conversion, model)
135	8.2	visitor_id → lead_id identity resolution	4	5	Conservative — never overwrites existing lead_id
136	8.4	/analytics/totals hero-card endpoint	1	—	Current + previous window with %-change chips; regex-pattern fix
137	5.2	LinkedIn UGC publisher	3	5	x-restli-id response header → canonical remote_id
138	5.3	X (Twitter) publisher	3	6	280-char ellipsis truncation; 403 maps to AuthError (write-scope)
139	5.4	Instagram Business publisher	3	6	Raises InstagramMissingMediaError — text-only impossible
140	11.3	GDPR org-export Celery task	3	6	Redact by DB column name (alias-aware); bytes always REDACTED
141	10.3	Workflow DAG runner	2	10	Defers on approval/branch — returns WorkflowRunResult.deferred_reason
142	11.1	/costs/totals + /costs/recent	2	—	Aggregates by kind + provider; previous-window delta

Files counts the modified+added Python and TypeScript files (excluding tests in this column; tests counted separately). "—" means PR added zero new test functions (typically frontend-only or pure plumbing). Total new tests in this sprint range: **81** across 15 new test files.

7.1 Distribution by phase

PRs landed per phase (depth pass)



22 depth-pass PRs distributed across 9 of the 12 phases (P0, P1, P2 had no new PRs this sprint — their primitives were complete pre-sprint).

CHAPTER 8

Architecture Snapshot

High-level shape of the system as of main @ 64242e8.

8.1 System diagram

Frontend · Next.js 14 App Router

44 pages · 26 <Dk*> components · Tailwind-on-tokens · light-mode only
/dashboard · /agent · /agents/{smm,seo,paid-media,analyst,creatives} · /inbox · /calendar · /channels · /integrations · /orgs · /admin/*

API · FastAPI

23 routers · org-membership-gated · 9 supervision roles
/agents, /approvals, /calendar, /channels, /integrations, /analytics, /costs, /email, /orgs/[id]/agent-threads, ...

Celery worker + beat

9 task modules · 9 beat schedules
scan_due_scheduled_posts (60s)
resolve_visitor_identities (05:30 UTC)
compute_daily_rollups (06:00 UTC)
compute_attribution (06:30 UTC)

Services layer · 12 modules

All provider-or-stub. 5 publishers in

publishers/. image_gen · asset_gen · email_send workflow_runner · mcp_client · quota_guard
publishers/{bluesky,linkedin,x,instagram} secret_box (Fernet) · embeddings · audit · storage Postgres
+ pgvector 21 models organizations, projects, leads, ... Redis Celery broker + backend MinIO / S3
Asset storage

8.2 Component inventory (snapshot)

COMPONENT CLASS	COUNT	NOTES
Backend Models (SQLAlchemy)	21	Organization, Project, BrandKit, Persona, Lead, Campaign, ScheduledPost, SocialAccount, Connection, ApprovalRequest, AuditEvent, AnalyticsEvent, Touchpoint, Conversion, AttributionResult, AnalyticsRollup, EmailTemplate, EmailCampaign, EmailSequence, EmailSequenceStep, Workflow, Playbook, TimeEntry, CostLedger, QuotaCounter, AgentThread, AgentMessage, ...

API routers (FastAPI)	23	One per domain area
Service modules	12	image_gen, asset_gen, email_send, workflow_runner, mcp_client, quota_guard, embeddings, secret_box, audit, mcp_registry, storage, ingestion, plus publishers/___init___
Publishers	5	publishers/{___init___, bluesky, linkedin, x, instagram}.py
Worker task modules	9	analytics, attribution, email_send, example, gdpr, identity, ingestion, publishing
Agent modules	10	_agent_base, anthropic_client, brand_style, conductor, creatives, smm, seo, paid_media, analyst, inbox
Frontend pages	44	Including dashboard, all five agent stations, calendar, inbox, channels, integrations, orgs, admin
<Dk*> components	26	Brand-vocabulary primitives
Backend test files	32	~221 individual test functions

CHAPTER 9

Test & Quality Gates

9.1 Test counts

FILE	TESTS	SPRINT?	NOTES
test_admin.py	7	—	pre-sprint
test_agents.py	10	—	pre-sprint — Creatives Agent
test_analytics.py	3	—	pre-sprint (legacy)
test_analytics_rollup.py	4	✓	Phase 8.1 daily rollups
test_approvals.py	9	—	pre-sprint
test_asset_gen.py	4	✓	Phase 3.3 video/voice/music
test_attribution_compute.py	7	✓	Phase 8.3 attribution allocation
test_auth.py	7	—	pre-sprint
test_bluesky_publisher.py	5	✓	Phase 5.1
test_brand_kits.py	8	—	pre-sprint
test_brand_style.py	5	✓	Phase 3.4 prompt composer
test_campaigns.py	7	—	pre-sprint
test_dashboard.py	1	—	pre-sprint
test_email_send.py	3	✓	Phase 7.1 Resend
test_gdpr_export.py	6	✓	Phase 11.3 GDPR export
test_goals.py	6	—	pre-sprint
test_health.py	1	—	pre-sprint
test_identity_resolution.py	5	✓	Phase 8.2
test_image_gen.py	4	✓	Phase 3.1 Replicate
test_ingestion.py	12	—	pre-sprint

test_instagram_publisher.py	6	✓	Phase 5.4
test_jobs.py	7	—	pre-sprint
test_kg.py	12	—	pre-sprint — pgvector KG
test_leads.py	9	—	pre-sprint
test_linkedin_publisher.py	5	✓	Phase 5.2
test_mcp_client.py	4	✓	Phase 6.1
test_orgs.py	8	—	pre-sprint
test_projects.py	7	—	pre-sprint
test_quota_guard.py	7	✓	Phase 11.2
test_scheduled_posts.py	5	—	pre-sprint (Phase 4)
test_storage.py	13	—	pre-sprint
test_workflow_runner.py	10	✓	Phase 10.3
test_x_publisher.py	6	✓	Phase 5.3
TOTAL	221	+81 SPRINT	15 NEW TEST FILES THIS SPRINT

9.2 CI status

GATE	STATUS	NOTES
backend-tests (pytest)	GREEN	All 221 tests pass on every PR; runs in ≈ 2m 30s
frontend-build (Next.js)	GREEN	~45s; production build with type-check
claude-review	INFORMATIONAL	Runs on every PR; <i>not</i> a required check (some PRs merged with this still pending)
poll-and-merge	PARTIALLY WIRED	Workflow exists but the auto-merge label doesn't actually invoke <code>gh pr merge --auto</code> — six PRs needed manual merging this sprint

9.3 What's NOT covered by tests yet

- End-to-end browser flows (Playwright) — none.
- Real provider integration tests (Replicate / Resend / Bluesky) — all stubs.
- Frontend component tests — none.
- Worker beat-task scheduling smoke tests.

CHAPTER 10

Risks, Tech Debt & Open Decisions

10.1 Active risks

RISK	L	I	MITIGATION
Cost ledger has no writers — /costs/totals reports \$0	H	M	Add cost-write hook on every adapter call; ½-day work
No real OAuth flow for any of 4 shipped publishers	H	H	OAuth lane in Sprint 2 (LinkedIn first; X needs paid tier; Instagram uses Meta Graph)
Auto-merge label doesn't auto-merge	M	L	GH Actions task to wire label → gh pr merge --auto
Workflow runner can't pause on approval	M	M	WorkflowRun model + paused state machine
Sandbox / dry-run mode missing — onboarding demos may post for real	M	H	Organization.sandbox_mode flag + adapter short-circuit
Sliding-window QuotaCounter has no writer — daily caps work, rolling don't	M	M	Quota writer in the same hook as cost-write
Anthropic provider rate-limited mid-sprint (one observed event)	M	L	Add retry-with-backoff in anthropic_client; consider Sonnet for lower-stakes calls
16 social platforms still on stub-only — channels page falsely advertises support	H	M	Gate channel cards on whether an adapter exists; or schedule per-platform PRs

L = Likelihood, I = Impact. H/M/L scale.

10.2 Tech debt

- **SocialAccount._interim_access_token** field name signals it's a placeholder — needs replacement with encrypted secret blob like Connection once OAuth ships.
- Existing v1 routers (campaigns, leads, analytics) are *not* Org-scoped yet (PR #117's TEMPORARY /api/v1/dashboard endpoint flags this).

- Pydantic class-based Config deprecation warnings on 14 schemas — should migrate to ConfigDict.
- poll-and-merge workflow runs but doesn't actually merge — leftover scaffold from earlier sprints.
- Frontend has both a `publisher_response.image_url` stash pattern (Instagram) and a TODO for `real asset_id` → URL resolution.

10.3 Open architecture decisions

1. **Claude Agent SDK vs raw Anthropic API.** Currently using raw API + `complete()` wrapper. Agent SDK gives tool-use loops + memory but locks us in. Decision needed before Phase 9.7 (programmatic Conductor delegation) work.
2. **WorkflowRun model shape.** Need to decide: one row per node execution? One row per run with JSON history? Affects audit story.
3. **Cost-write contract.** Do adapters write CostLedger directly, or emit an event that a worker consumes? Direct writes are simpler; events are more replay-friendly.
4. **Where do real Claude Agent SDK tool calls register?** Need to bridge Agent SDK tool list ↔ MCP tool list ↔ approval-required-action list.
5. **Sandbox semantics.** Strict — every external call short-circuits to a stub? Or soft — calls succeed but flag the resulting row? Strict is simpler; soft is more realistic.

CHAPTER 11


Sprint 2 — Recommended Roadmap

If Sprint 2 mirrors Sprint 1 in length (~3 days), here is the ranked work plan to push the platform from $\approx 52\%$ complete to $\approx 75\%$ complete.

11.1 Lane priorities

#	LANE	WHY IT'S FIRST	EST. EFFORT	PHASE MOVED
1	Cost-write hooks on every adapter	Makes /costs/totals real. Highest ROI in the whole plan — one-liner per adapter.	½ day	11: 40 → 55%
2	Real OAuth flow — LinkedIn first	Unblocks the four shipped publishers for production. LinkedIn is the cleanest OAuth 2.0 grant.	1 day	5: 30 → 50%
3	Brand Setup Studio UI	Closes the biggest backend-vs-frontend delta in the plan. Backend 100%, UI 0%.	1 day	2: 50 → 75%
4	Sequence builder (react-flow)	Biggest single UI gap on Phase 7; unlocks email automation product.	1.5 days	7: 30 → 55%
5	WorkflowRun + HTTP runner endpoint	Turns the new runner from library to product; unblocks approval-pause and branch nodes.	½ day	10: 20 → 35%
6	Sandbox / dry-run mode	Closes the highest-impact compliance gap; needed for demo safety.	½ day	11: 55 → 70%
7	Stripe invoicing	Largest revenue-side feature for agencies. New model + migration.	1.5 days	10: 35 → 55%
8	HubSpot CRM 2-way sync via MCP	Best-known CRM; brings Phase 8 from analytics-only to full ops.	1.5 days	8: 50 → 70%

11.2 Projected plan completeness after Sprint 2



Light purple is current. Darker purple is Sprint 2 projection. Net plan completeness target: $\approx 75\%$.

11.3 Out of scope for Sprint 2

Explicitly punted (the user has signalled to handle these themselves, or they're nice-to-have):

- Marketing assets (#52 demo video, #53 launch posts) — user has these.
- Helm chart rebuild (#24) — deployment polish, not blocking demos.
- Time-decay and Markov attribution models — interesting but not on the critical path.
- 13 of the 20 social platform adapters — covered after Sprint 2 in a "long-tail" sprint.

APPENDIX A

Closed Issues with PR Mapping

Issues that moved to closed during this sprint, with the PR that closed them.

ISSUE	TITLE	CLOSED BY
#31	[agents] Claude Agent SDK integration scaffold	foundation work — pre-sprint
#32	[agents] Conductor agent shell	#121
#33	[agents] Creatives Agent (end-to-end MVP)	#72 (pre-sprint), reinforced by #128
#34	[agents] Approval Inbox UI	#73 (pre-sprint), reinforced by #126
#35	[agents] SMM Agent + Calendar Station	#122
#36	[agents] SEO Agent + Search Station	#122
#37	[agents] Paid Media Agent + Spend Station	#122
#38	[agents] Analyst Agent + Insights Station	#122
#39	[publishing] SocialAccount model	pre-sprint (PR #115)
#43	[publishing] Email sender via Resend (real)	#127
#44–#48	Docs / README / changelog / API ref / architecture	pre-sprint
#49–#51	Marketing one-pager / slides / video script	pre-sprint
#57	[release] Tag v0.1.0-mvp	pre-sprint
#78	[bug] pgvector startup crash	pre-sprint #79
#80	[bug] colima IPv6 frontend fetch	pre-sprint #81
#82	[bug] EmailStr .local TLD reject	pre-sprint #83
#88	[Epic] Phase 3 — Content Generation Pipeline	#124 + #125 + #128
#91	[Epic] Phase 6 — MCP Integration Hub	#130
#92	[Epic] Phase 7 — Email + Ads + Sequences	#127 + #131 + #133
#93	[Epic] Phase 8 — Lead 2.0 + CRM + Attribution	#132 + #134 + #135 + #136
#94	[Epic] Phase 9 — Agent Fleet	#121 + #122 + #123

APPENDIX B

Open Issues — Status & Owners

ISSUE	TITLE	STATUS	NOTE
#85	Phase 0 Epic	Stale	Stories #97–#104 mostly delivered; close as hygiene
#86	Phase 1 Epic	Open	Tabbed /orgs/[id] + member invite UI
#87	Phase 2 Epic	Open	Brand Setup Studio, KG Console
#89	Phase 4 Epic	Open	Conflict detection + best-time recommender
#90	Phase 5 Epic	Open	OAuth flows + 16 more platforms
#95	Phase 11 Epic	Open	GDPR done (#140), costs endpoint done (#142), quota guard done (#133). Remaining: sandbox, client portal, cost-write hooks, sliding-window quota writer
#96	Phase 10 Epic	Open	Workflow runner done (#141). Remaining: onboarding wizard, Kanban, time tracker UI, Stripe invoicing, PDF reports
#97–#104	Phase 0 design stories	Delivered, not closed	One-time cleanup task
#24	Helm chart rebuild	Open	Deferred — not on critical path
#28	Q4 Freshness re-ingestion	Open	Future
#30	Q6 Project Setup Wizard	Open	Sprint 3
#40	X publisher mock	Delivered, not closed	Closed by #138 — close manually
#41	LinkedIn publisher mock	Delivered, not closed	Closed by #137 — close manually
#42	Instagram publisher mock	Delivered, not closed	Closed by #139 — close manually
#52	Demo video recording	Open · user-owned	User has signalled they'll handle marketing assets

#53	Launch social posts (drafts)	Open · user-owned	Same
-----	------------------------------	-------------------	------

APPENDIX C

Full PR Catalogue Table

All 22 merged PRs in the depth-pass range, plus the one in flight at snapshot time.

#	DATE	PHASE	TITLE	BRANCH (SQUASHED)
121	05-12	9.1	Conductor real Claude	feat/9.1-conductor-real-claude
122	05-12	9.2	SMM/SEO/Paid/Analyst real Claude	feat/9.2-smm-agent
123	05-12	9.3	Shared DkAgentChat	feat/9.3-role-agent-chat-ui
124	05-12	3.1	Replicate image-gen adapter	feat/3.1-image-gen-adapter
125	05-12	3.3	Video/voice/music adapters	feat/3.3-video-voice-music
126	05-12	4.2	Multimodal inbox previews	feat/4.2-inbox-multimodal-previews
127	05-12	7.1	Resend send adapter	feat/7.1-email-send-adapter
128	05-12	3.4	Brand-aware visual prompts	feat/3.4-brand-aware-image-prompts
129	05-12	5.1	Bluesky publisher	feat/5.1-bluesky-publisher
130	05-12	6.1	Async MCP client	feat/6.1-mcp-async-client
131	05-12	7.2	Approve → email delivery	feat/7.2-approval-triggers-email-send
132	05-12	8.1	Daily analytics rollup	feat/8.1-analytics-daily-rollup
133	05-12	7.3 + 11.2	Campaign hard-gate + quota guard	feat/7.3-email-campaign-send

134	05-12	8.3	Attribution compute	feat/8.3-attribution-compute
135	05-12	8.2	Identity resolution	feat/8.2-identity-resolution
136	05-12	8.4	/analytics/totals	feat/8.4-analytics-totals
137	05-12	5.2	LinkedIn publisher	feat/5.2-linkedin-publisher
138	05-12	5.3	X publisher	feat/5.3-x-twitter-publisher
139	05-12	5.4	Instagram publisher	feat/5.4-instagram-publisher
140	05-12	11.3	GDPR export task	feat/11.3-gdpr-export
141	05-12	10.3	Workflow DAG runner	feat/10.3-workflow-runner
142	05-12	11.1	/costs endpoints	feat/11.1-cost-ledger-totals

APPENDIX D

Glossary & Acronyms

TERM	MEANING
Approval Inbox	The Hard-gate decision queue. Every outbound action that the org's autonomy posture doesn't auto-approve lands here for human review.
BrandKit	Versioned brand identity per Org — palette, fonts, voice sliders, do-say / don't-say, positioning. Agents read from the active one.
Conductor	The Manager-tier agent (<code>AgentKind.conductor</code>). Decomposes briefs into work for role-agents.
CostLedger	Per-call provider spend table. Daily rollups read from it.
DkAgentChat	The shared chat component on every role-agent station, parameterised by <code>AgentKind</code> .
Hard-gate	The default supervision posture: agents never publish; everything routes through Approval. PLAN-v1.2 §v2.0 §5.2.
Identity resolution	The 05:30 UTC beat job that back-stamps <code>lead_id</code> on anonymous Touchpoints sharing a <code>visitor_id</code> .
MCP	Model Context Protocol — the standard for letting an LLM call an external tool. We use a 35-server registry + an async HTTP client.
Org / Project	Multi-tenancy hierarchy. Every domain row has both IDs; every API call is scoped to memberships.
Role agent	One of the four specialists: <code>smm</code> (Social Media Manager), <code>seo</code> , <code>paid_media</code> , <code>analyst</code> .
Supervision posture	Per-action-class setting — <code>hard_gate</code> , <code>soft_gate</code> , <code>auto</code> . Stored on <code>Organization.autonomy_posture_json</code> .
Workflow DSL	The JSON shape of <code>Workflow.dsl_json</code> : nodes + edges. v1 supports <code>llm</code> , <code>tool_call</code> , <code>noop</code> node types.

END OF REPORT

**Three days.
Twenty-three pull
requests.
One agent-driven
marketing OS.**

Dev Sprint 1 Status Report · v1.0

Generated 2026-05-12 21:00 IST

Snapshot ref · main @ 64242e8

Engineering · with Claude Code (Opus 4.7, 1M context)

Confidential — internal distribution · DClaw Marketing